

## Hive Assignment-2

### Hive Practical questions:

#### Hive Join operations

Create a table named CUSTOMERS(ID | NAME | AGE | ADDRESS | SALARY)

```
hive> create table customers(id int ,name string ,age int ,address string ,salary double)
> row format delimited
> fields terminated by ',';
```

customers.id	customers.name	customers.age	customers.address	customers.salary
1	Nimeshika	21	Madurai	22000.0
2	Sowmya	23	Coimbatore	25000.0
3	Madhesh	22	Virudhunagar	23000.0
4	Vijay	24	Villupuram	50000.0
5	Kiran	26	Bangalore	29000.0
6	Maruthu	24	Ramnad	27000.0

Create a Second table ORDER(OID | DATE | CUSTOMER\_ID | AMOUNT)

```
hive> create table orders(oid int, order_date string,customer_id int ,amount double)
> row format delimited
> fields terminated by ',';
```

orders.oid	orders.order_date	orders.customer_id	orders.amount
1	29-03-2023	2	50.0
2	30/23/2023	4	40.0
3	01-01-2023	6	20.0

Now perform different joins operations on top of these tables

(Inner JOIN, LEFT OUTER JOIN ,RIGHT OUTER JOIN ,FULL OUTER JOIN)

#### Inner Join

```
hive> select c.id,c.name,c.address,o.amount from orders o
> inner join customers c
> on o.customer_id = c.id;
```

c.id	c.name	c.address	o.amount
2	Sowmya	Coimbatore	50.0
4	Vijay	Villupuram	40.0
6	Maruthu	Ramnad	20.0

## Left Join

```
hive> select c.id,c.name,c.address,o.amount from orders o
> left join customers c
> on o.customer_id = c.id
> where o.amount>40;
```

OK

c.id	c.name	c.address	o.amount
2	Sowmya	Coimbatore	50.0

## Right Join

```
hive> select c.id,c.name,c.address,o.amount from orders o
> right join customers c
> on o.customer_id = c.id;
```

c.id	c.name	c.address	o.amount
1	Nimeshika	Madurai	NULL
2	Sowmya	Coimbatore	50.0
3	Madhesh	Virudhunagar	NULL
4	Vijay	Villupuram	40.0
5	Kiran	Bangalore	NULL
6	Maruthu	Ramnad	20.0

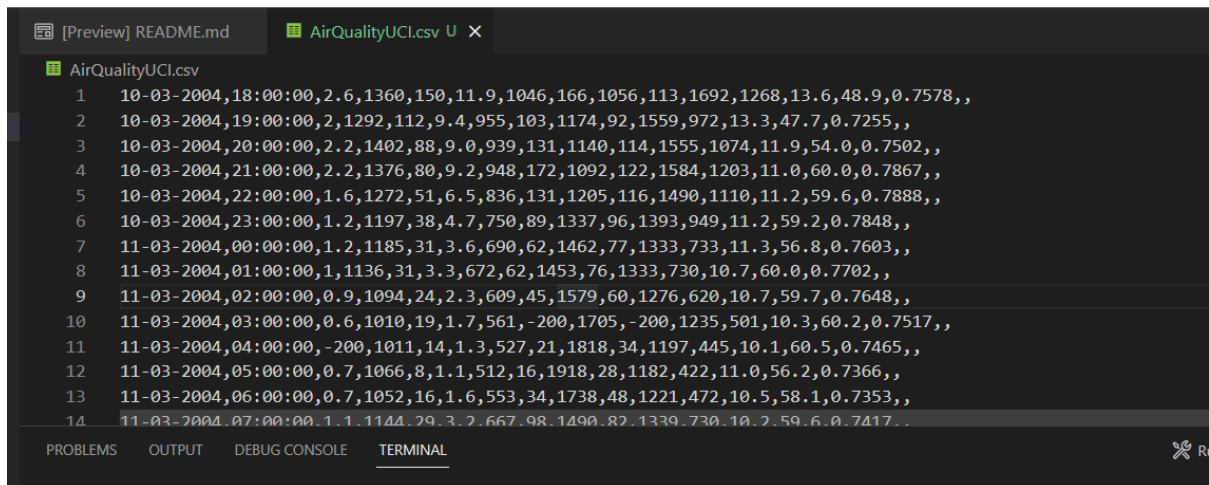
## Full Join

```
hive> select c.id,c.name,c.address,o.amount from orders o
> full join customers c
> on o.customer_id = c.id;
```

c.id	c.name	c.address	o.amount
1	Nimeshika	Madurai	NULL
2	Sowmya	Coimbatore	50.0
3	Madhesh	Virudhunagar	NULL
4	Vijay	Villupuram	40.0
5	Kiran	Bangalore	NULL
6	Maruthu	Ramnad	20.0

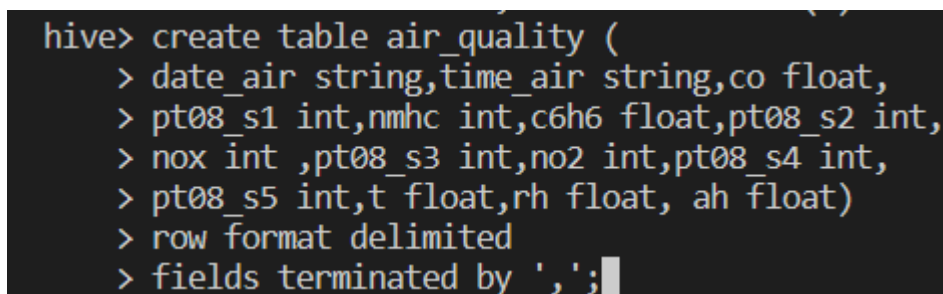
Download a data from the given location -

<https://archive.ics.uci.edu/ml/machine-learning-databases/00360/>



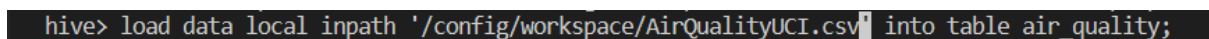
The screenshot shows a code editor with a file named 'AirQualityUCI.csv' open. The file contains a list of 14 rows of data, each representing an air quality measurement. The data is comma-separated and includes timestamps, various pollutant concentrations (like CO, NO2, PM10, etc.), and other environmental factors. The editor has a dark theme and a terminal window at the bottom.

1. Create a hive table as per given schema in your dataset



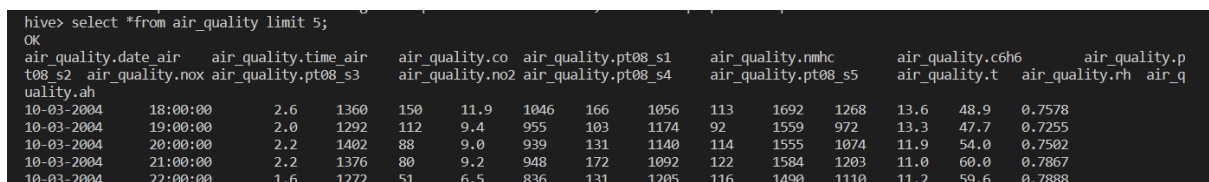
```
hive> create table air_quality (  
  > date_air string,time_air string,co float,  
  > pt08_s1 int,nmhc int,c6h6 float,pt08_s2 int,  
  > nox int ,pt08_s3 int,no2 int,pt08_s4 int,  
  > pt08_s5 int,t float,rh float, ah float)  
  > row format delimited  
  > fields terminated by ',';
```

2. try to place a data into table location



```
hive> load data local inpath '/config/workspace/AirQualityUCI.csv' into table air_quality;
```

3. Perform a select operation .



```
hive> select *from air_quality limit 5;  
OK  
air_quality.date_air  air_quality.time_air  air_quality.co  air_quality.pt08_s1  air_quality.nmhc  air_quality.c6h6  air_quality.p  
t08_s2  air_quality.nox  air_quality.pt08_s3  air_quality.no2  air_quality.pt08_s4  air_quality.pt08_s5  air_quality.t  air_quality.rh  air_q  
uality.ah  
10-03-2004  18:00:00  2.6  1360  150  11.9  1046  166  1056  113  1692  1268  13.6  48.9  0.7578  
10-03-2004  19:00:00  2.0  1292  112  9.4  955  103  1174  92  1559  972  13.3  47.7  0.7255  
10-03-2004  20:00:00  2.2  1402  88  9.0  939  131  1140  114  1555  1074  11.9  54.0  0.7502  
10-03-2004  21:00:00  2.2  1376  80  9.2  948  172  1092  122  1584  1203  11.0  60.0  0.7867  
10-03-2004  22:00:00  1.6  1272  51  6.5  836  131  1205  116  1490  1110  11.2  59.6  0.7888
```

```
hive> select date_air,ah,nmhc from air_quality limit 10;
OK
date_air      ah      nmhc
10-03-2004    0.7578  150
10-03-2004    0.7255  112
10-03-2004    0.7502  88
10-03-2004    0.7867  80
10-03-2004    0.7888  51
10-03-2004    0.7848  38
11-03-2004    0.7603  31
11-03-2004    0.7702  31
11-03-2004    0.7648  24
11-03-2004    0.7517  19
Time taken: 0.162 seconds, Fetched: 10 row(s)
```

4. Fetch the result of the select operation in your local as a csv file .

```
hive> INSERT OVERWRITE LOCAL DIRECTORY '/config/workspace/Output_air_quality'
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> select date_air,ah,nmhc from air_quality limit 10;
```

The screenshot shows the VS Code interface. On the left, the Explorer panel displays the workspace structure with a folder named 'Output\_air\_quality' containing a file '000000\_0'. The main editor area shows the content of 'Output\_air\_quality.csv', which contains 11 rows of data (rows 1-11) in CSV format, matching the output of the Hive query shown in the previous block.

Line	date_air	ah	nmhc
1	11-03-2004	0.7517	19
2	11-03-2004	0.7648	24
3	11-03-2004	0.7702	31
4	11-03-2004	0.7603	31
5	10-03-2004	0.7848	38
6	10-03-2004	0.7888	51
7	10-03-2004	0.7867	80
8	10-03-2004	0.7502	88
9	10-03-2004	0.7255	112
10	10-03-2004	0.7578	150
11			

5. Perform group by operation.

```
hive> select date_air,sum(pt08_s1) from air_quality
> group by date_air;
```

Total MapReduce CPU Time		
OK		
NULL		
01-01-2005		26742
01-02-2005		30004
01-03-2005		19684
01-04-2004		25532
01-04-2005		21679
01-05-2004		26340
01-06-2004		27254
01-07-2004		27134
01-08-2004		23380
01-09-2004		25454
01-10-2004		30061
01-11-2004		29844
01-12-2004		28161
02-01-2005		22364
02-02-2005		28254
02-03-2005		23302
02-04-2004		30934

7. Perform filter operation at least 5 kinds of filter examples .

```
hive> select date_air,pt08_s1,pt08_s2,pt08_s3 from air_quality where date_air >'01-01-2004' and date_air <'31-12-2004' limit 50;
OK
```

OK			
10-03-2004	1360	1046	1056
10-03-2004	1292	955	1174
10-03-2004	1402	939	1140
10-03-2004	1376	948	1092
10-03-2004	1272	836	1205
10-03-2004	1197	750	1337
11-03-2004	1185	690	1462
11-03-2004	1136	672	1453
11-03-2004	1094	609	1579
11-03-2004	1010	561	1705
11-03-2004	1011	527	1818
11-03-2004	1066	512	1918
11-03-2004	1052	553	1738
11-03-2004	1144	667	1490

```
hive> select pt08_s1 from air_quality where pt08_s1>2000;
```

```
pt08_s1
2040
2008
```

```
hive> select date_air,pt08_s1,pt08_s2 from air_quality where pt08_s1>2000;
OK
date_air      pt08_s1 pt08_s2
15-03-2004    2040    1754
23-11-2004    2008    1980
```

```
hive> select date_air,co from air_quality where co>7 limit 5;
OK
date_air      co
15-03-2004    8.1
15-03-2004    8.0
17-03-2004    7.6
15-04-2004    7.3
29-04-2004    7.2
```

```
hive> select date_air,c6h6 from air_quality where c6h6>6.0 and c6h6<7.0 limit 5;
OK
date_air      c6h6
10-03-2004    6.5
11-03-2004    6.3
12-03-2004    6.4
13-03-2004    6.2
14-03-2004    6.4
Time taken: 0.18 seconds, Fetched: 5 row(s)
```

## 8. show and example of regex operation

```
hive> select date_air,c6h6 from air_quality where date_air regexp '^01';
OK
01-04-2004    6.3
01-04-2004    5.1
01-04-2004    4.1
01-04-2004    4.0
01-04-2004    2.4
01-04-2004    2.9
01-04-2004    7.4
01-04-2004    19.8
01-04-2004    31.9
01-04-2004    21.0
01-04-2004    11.7
01-04-2004    6.8
01-04-2004    7.4
01-04-2004    8.6
01-04-2004    -200.0
01-04-2004    -200.0
01-04-2004    -200.0
01-04-2004    16.4
```

## 9. alter table operation

```
FAILED: ParseException line 1:28 cannot recognize input near 'add column' '(' in alter table
hive> alter table air_quality add columns (temperature string comment 'temperature of air');
OK
Time taken: 0.208 seconds, Fetched: 3 row(s)
```

air_quality.date	air	air_quality.time	air	air_quality.co	air_quality.pt08_s1	air_quality.nmhc	air_quality.c6h6	air_quality.p						
t08_s2	air_quality.no2	air_quality.pt08_s3	air_quality.no2	air_quality.pt08_s4	air_quality.pt08_s5	air_quality.t	air_quality.rh	air_q						
uality.ah	air_quality.temperature													
10-03-2004	18:00:00	2.6	1360	150	11.9	1046	166	1056	113	1692	1268	13.6	48.9	0.7578
10-03-2004	19:00:00	2.0	1292	112	9.4	955	103	1174	92	1559	972	13.3	47.7	0.7255
10-03-2004	20:00:00	2.2	1402	88	9.0	939	131	1140	114	1555	1074	11.9	54.0	0.7502

Time taken: 0.208 seconds, Fetched: 3 row(s)

## 10 . drop table operation

```
Time taken: 0.174 seconds, Fetched: 10 row(s)
hive> drop table customers;
OK
Time taken: 0.484 seconds
hive> select *from customers;
FAILED: SemanticException [Error 10001]: Line 1:13 Table not found 'customers'
hive>
```

## 12 . order by operation .

```
hive> select date_air,pt08_s1 from air_quality
> order by pt08_s1 desc limit 50;
```

```
OK
date_air      pt08_s1
15-03-2004    2040
23-11-2004    2008
23-11-2004    1982
17-03-2004    1975
17-03-2004    1973
15-03-2004    1961
26-11-2004    1956
18-03-2004    1934
23-11-2004    1918
15-03-2004    1917
04-11-2004    1915
26-10-2004    1908
14-03-2004    1898
15-03-2004    1895
20-10-2004    1884
04-11-2004    1882
13-12-2004    1881
15-04-2004    1875
05-11-2004    1870
```

13 . where clause operations you have to perform .

```
Time taken: 0.249 seconds, Fetched: 2 row(s)
hive> select pt08_s1 from air_quality where pt08_s1>2000;
OK
pt08_s1
2040
2008
Time taken: 0.249 seconds, Fetched: 2 row(s)
```

14 . sorting operation you have to perform .

```
Time taken: 0.249 seconds, Fetched: 2 row(s)
hive> select air_date,pt08_s3 from air_quality sort by pt08_s3 asc;
```

22-03-2004	1417
30-05-2004	1417
17-03-2004	1417
27-05-2004	1417
30-03-2004	1418
10-11-2004	1419
13-10-2004	1419
13-03-2004	1420
03-06-2004	1420
31-05-2004	1421
24-01-2005	1421
27-04-2004	1423
27-04-2004	1424
26-12-2004	1425
09-12-2004	1426
13-06-2004	1426
20-04-2004	1429
10-11-2004	1433
06-03-2005	1433
14-10-2004	1434
25-09-2004	1434

15 . distinct operation you have to perform .

```
Time taken: 24.500 seconds, Fetched: 3471 row(s)
hive> select distinct(date_air) from air_quality;
Query ID = abc_20220222170542_1a2ab424_f817_4020...
```



```
OK
date_air
01-01-2005
01-02-2005
01-03-2005
01-04-2004
01-04-2005
01-05-2004
01-06-2004
01-07-2004
01-08-2004
01-09-2004
01-10-2004
01-11-2004
01-12-2004
02-01-2005
02-02-2005
02-03-2005
```

16 . like an operation you have to perform .

```
hive> select date_air from air_quality where date_air like '^01';
OK
```

[illegible]

17 . union operation you have to perform .

```
hive> select date_air from air_quality  
> union  
> select name from customers;
```

```
Query ID = hls_20220222171420_ok702aaa_2  
30-09-2004  
30-10-2004  
30-11-2004  
30-12-2004  
31-01-2005  
31-03-2004  
31-03-2005  
31-05-2004  
31-07-2004  
31-08-2004  
31-10-2004  
31-12-2004  
Date  
Kiran  
Madhesh  
Maruthu  
Nimeshika  
Sowmya  
Vijay  
Time taken: 25.278 seconds, Fetched: 399 row(s)  
hive>
```

18 . table view operation you have to perform.

```
Time taken: 0.177 seconds, Fetched: 10 row(s)  
hive> create view air_view as  
> select pt08_s1,pt08_s2 from air_quality;  
OK  
pt08_s1 pt08_s2
```

```
air_data_view.pt08_s1  air_data_view.pt08_s2
NULL      NULL
1360      1046
1292      955
1402      939
1376      948
1272      836
1197      750
1185      690
1136      672
1094      609
```

## **Scenario Based questions:**

**1. Will the reducer work or not if you use "Limit 1" in any HiveQL query?**

Reducer will not work if we use limit in select clause,

When used in the other aggregation query even though if we use limit in the query the reducer will work.

**2. Suppose I have installed Apache Hive on top of my Hadoop cluster using default meta store configuration. Then, what will happen if we have multiple clients trying to access Hive at the same time?**

The default meta store configuration allows only one hive session to be opened at a time for accessing the metastore.

If multiple clients try to access the metastore then they will get an error.

**3. Suppose, I create a table that contains details of all the transactions done by the customers: CREATE TABLE transaction\_details (cust\_id INT, amount FLOAT, month STRING, country STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' ;**

**Now, after inserting 50,000 records in this table, I want to know the total revenue generated for each month. But, Hive is taking too much time in processing this query. How will you solve this problem and list the steps that I will be taking in order to do so?**

We can solve this issue by Partitioning the table according to month.

Hence we will be scanning the data for particular month and not the whole dataset.

Steps are:

a)create a partition table

```
CREATE TABLE transaction_details_partition  
(cust_id INT, amount FLOAT, month STRING, country  
STRING)  
  
PARTITIONED BY (month STRING)  
  
ROW FORMAT DELIMITED  
  
FIELDS TERMINATED BY ',' ;
```

b)enable dynamic partitioning

```
SET hive.exec.dynamic.partition = true;
```

c)transfer the data from non partitioned table to partitioned table.

```
INSERT OVERWRITE TABLE transaction_details_partition  
(month) SELECT cust_id, amount, country, month FROM  
transaction_details;
```

**4.How can you add a new partition for the month December in the above partitioned table?**

With the help of alter table command

```
ALTER TABLE transaction_details_partition ADD PARTITION  
(month='December') location "
```

**5. I am inserting data into a table based on partitions dynamically. But, I received an error – FAILED ERROR IN SEMANTIC ANALYSIS: Dynamic partition strict mode requires at least one static partition column. How will you remove this error?**

To remove this error we need to execute the commands as,

```
SET hive.exec.dynamic.partition=true;
```

```
SET hive.exec.dynamic.partition.mode=nonstrict;
```

**6. Suppose, I have a CSV file – ‘sample.csv’ present in ‘/temp’ directory with the following entries:**

**id first\_name last\_name email gender ip\_address**

**How will you consume this CSV file into the Hive warehouse using built-in SerDe?**

SERDE allows us to convert the unstructured bytes into records then we can process it using hive.

```
CREATE EXTERNAL TABLE sample
```

```
(id int, first_name string,
```

```
last_name string, email string,
```

```
gender string, ip_address string)
```

```
ROW FORMAT SERDE
```

```
‘org.apache.hadoop.hive.serde2.OpenCSVSerde’
```

```
STORED AS TEXTFILE LOCATION “;
```

**7. Suppose, I have a lot of small CSV files present in the input directory in HDFS and I want to create a single Hive table corresponding to these files. The data in these files are in the format: {id, name, e-mail, country}. Now, as we know, Hadoop performance degrades when we use lots of small files.**

**So, how will you solve this problem where we want to create a single Hive table for lots of small files without degrading the performance of the system?**

It can be resolved using SequenceFile(it will help to group the small files into a single file)

Create a temporary table:

```
CREATE TABLE temp_table (id INT, name STRING, e-mail STRING,  
country STRING)
```

ROW FORMAT FIELDS DELIMITED TERMINATED BY ',' STORED AS TEXTFILE;

Load the data into temp\_table:

LOAD DATA INPATH '/input' INTO TABLE temp\_table;

Create a table that will store data in SequenceFile format:

CREATE TABLE sample\_seqfile (id INT, name STRING, e-mail STRING, country STRING)

ROW FORMAT FIELDS DELIMITED TERMINATED BY ',' STORED AS SEQUENCEFILE;

Transfer the data from the temporary table into the sample\_seqfile table:

INSERT OVERWRITE TABLE sample SELECT \* FROM temp\_table;

**8.LOAD DATA LOCAL INPATH 'Home/country/state/'**

**OVERWRITE INTO TABLE address;**

**The following statement failed to execute. What can be the cause?**

Several errors may occur as,

- a)Syntax error
- b)Permission issue
- c)Missing file
- d)Table not found
- e)Overwrite issue

**9.Is it possible to add 100 nodes when we already have 100 nodes in Hive? If yes, how?**

Yes, it is possible to add 100 nodes to an existing Hive cluster with 100 nodes. The process of adding more nodes to an existing Hive cluster is known as scaling out.

Here are the general steps to add more nodes to a Hive cluster:

- a) Install the necessary software on the new nodes, such as Hadoop and Hive.
- b) Configure the new nodes to join the existing cluster. This involves updating the configuration files of the new nodes with the information about the existing nodes and their roles in the cluster.
- c) Start the Hadoop and Hive services on the new nodes.
- d) Once the new nodes have joined the cluster and are up and running, you can rebalance the data across the nodes to ensure that the workload is distributed evenly.